

2024-07-09

OpenHPC: Beyond the Install Guide

OpenHPC: Beyond the Install Guide
for PEARC24

Sharon Colson Jim Moroney Mike Renfro

Tennessee Tech University

2024-07-22

2024-07-09

OpenHPC: Beyond the Install Guide

└─ Introduction

└─ Acknowledgments and shameless plugs

└─ Acknowledgments and shameless plugs

x

Acknowledgments and shameless plugs

[OpenHPC](#) especially Tim Middelkoop (Internet2) and Chris Simmons (Massachusetts Green High Performance Computing Center). They have a BOF at 1:30 Wednesday. You should go to it.

[Jetstream2](#) especially Jeremy Fischer, Mike Lowe, and Julian Pistorius. Jetstream2 has a tutorial at the same time as this one. Please stay here.

[NSF CC*](#) for the equipment that led to some of the lessons we're sharing today (award #2127188).

[ACCESS](#) current maintainers of the project formerly known as the XSEDE Compatible Basic Cluster.

2024-07-09

OpenHPC: Beyond the Install Guide

└ Introduction

└ Where we're starting from

└ Where we're starting from

x

Where we're starting from



Figure 1: Two example HPC networks

31 HPC clusters (2 shown) with:

1. Rocky Linux 9
2. OpenHPC 3
3. Warewulf 3
4. Slurm
5. 2 non-GPU nodes
6. 2 GPU nodes (currently without GPU drivers, so: expensive non-GPU nodes)
7. 1 management node (SMS)
8. 1 unprovisioned login node

OpenHPC: Beyond the Install Guide

└─ Introduction

└─ Where we're starting from

└─ Where we're starting from

Where we're starting from

We used the OpenHPC automatic installation script from [Appendix A](#) with a few variations:

1. Installed `a-mail` to have a valid MailProg for `slurm.conf`.
2. Created `user1` and `user2` accounts with password-less sudo privileges.
3. Changed `CHROOT` from `/opt/ohpc/admin/images/rocky0.3` to `/opt/ohpc/admin/images/rocky0.4`.
4. Enabled `slurmd` and `slunge` in `CHROOT`.
5. Added `nano` and `yum` to `CHROOT`.
6. Removed a redundant `ReturnToService` line from `/etc/slurm/slurm.conf`.
7. Stored all nodes' SSH host keys in `/etc/ssh/ssh_known_hosts`.

x

2024-07-09

OpenHPC: Beyond the Install Guide

└─ Introduction

└─ Where we're going

└─ Where we're going

Where we're going

1. A slightly more secured SMS
2. A login node that's practically identical to a compute node (except for where it needs to be different)
3. GPU drivers on the GPU nodes
4. Using node-local storage for the OS and/or scratch
5. De-coupling the SMS and the compute nodes (e.g., independent kernel versions)
6. Easier management of node differences (GPU or not, diskless/single-disk/multi-disk, InfiniBand or not, etc.)
7. Slurm configuration to match some common policy goals (fair share, resource limits, etc.)

x

OpenHPC: Beyond the Install Guide

└─ Configuring Slurm policies

└─ Sample slide

This is my note.

- It can contain Markdown
- like this list

Left column

This slide has two columns. They don't always have to have columns. It also has a titled block of content in the left column. Make sure you've always got a `::: notes` block after the slide content, even if it has no content.

Use `#` and `##` headers in the Markdown file to make level-1 and level-2 headings. `###` headers to make slide titles, and `####` to make block titles.